

cycle-gan-p100

April 29, 2024

1 Image to Monet Cycle GAN

This project aims to solve this [Kaggle Contest](#). I based my code in the following [Cycle GAN Implementation Guide](#).

1.1 Objective

The main objective of this contest is to generate Monet images from normal images. Submit your generated images and be top target in the leaderboard.

1.2 Architecture

As the contest ask you the submit format needs to be a notebook, I am using the GPU P100 architecture.

1.3 Initial imports

I will need some libraries for data processing and PyTorch for easy to use training a Cycle GAN.

```
[ ]: import cv2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os

import torch
import torch.nn as nn
import torch.optim as optim

from tqdm import tqdm
from collections import defaultdict

import random
from torch.utils.data import Dataset, DataLoader
```

1.4 Used Folders listing

```
[ ]: root_path = "/kaggle/input/gan-getting-started"
os.listdir(root_path)
```

1.5 Data Exploration

First, I want to see what type of image I am dealing with, the size and ways of how to read it and processes it.

```
[ ]: read_img = lambda path: cv2.cvtColor(cv2.imread(path), cv2.COLOR_BGR2RGB)
```

```
[ ]: data_path = f"{root_path}/photo_jpg"
sample_photo = read_img(os.path.join(data_path, os.listdir(data_path)[0]))

data_path = f"{root_path}/monet_jpg"
sample_monet = read_img(os.path.join(data_path, os.listdir(data_path)[0]))
```

```
[ ]: sample_photo.shape
```

```
[ ]: sample_photo.min(), sample_photo.max(), sample_photo.dtype
```

```
[ ]: plt.subplot(121)
plt.title("Photo")
plt.imshow(sample_photo)

plt.subplot(122)
plt.title("Photo")
plt.imshow(sample_monet)
```

1.6 Cycle GAN Implementation

A cycle GAN is similar to the traditional GAN architecture, but it do not need equivalent pairs of images to learn the context of the desired output. Such as the example above shows, it not exist equivalent pairs in this dataset.

1.6.1 Initialization

The size of the images are 256x256 images with 3 channels, I will save it in the following global variables.

```
[ ]: IMG_H, IMG_W, IMG_C = 256, 256, 3
```

And then starting the implementation...

1.6.2 Downsampling

The first part of a traditional Cycle GAN architecture is the downsampling, it reduces the resolution of the images giving the network the most relevant information of the image.

```
[ ]: def downsample(in_channels, out_channels, kernel_size, norm=True):
    downsample_block = [nn.Conv2d(in_channels=in_channels,
                                   out_channels=out_channels,
                                   kernel_size=kernel_size,
                                   stride=2,
                                   padding=(kernel_size - 1) // 2,
                                   bias=False)

    ]

    if norm:
        downsample_block.append(nn.GroupNorm(num_groups=out_channels,
                                              num_channels=out_channels))

    downsample_block.append(nn.LeakyReLU())
    return nn.Sequential(*downsample_block)
```

```
[ ]: x = torch.randn(1, IMG_C, IMG_H, IMG_W)

out = downsample(3, 3, 3)(x)

assert IMG_C == out.shape[1]
assert IMG_H == out.shape[2] * 2
assert IMG_W == out.shape[3] * 2
```

1.6.3 Upsampling

You can define the upsampling step as the opposite of downsampling. From information, it upscales the image generating a new one.

```
[ ]: def upsample(in_channels, out_channels, kernel_size, dropout=False):
    downsample_block = [nn.ConvTranspose2d(in_channels=in_channels,
                                           out_channels=out_channels,
                                           kernel_size=kernel_size,
                                           stride=2,
                                           padding=(kernel_size - 1) // 2,
                                           bias=False),

    nn.GroupNorm(num_groups=out_channels,
                num_channels=out_channels)

    ]

    if dropout:
        downsample_block.append(nn.Dropout(0.5))

    downsample_block.append(nn.ReLU())
    return nn.Sequential(*downsample_block)
```

```
[ ]: x = torch.randn(1, IMG_C, IMG_H, IMG_W)

out = upsample(3, 3, 4)(x)
```

```

assert IMG_C == out.shape[1]
assert IMG_H == out.shape[2] // 2
assert IMG_W == out.shape[3] // 2

```

Usually, it is needed to initialize the network convolutional layers with random values, the following cell do this.

```

[ ]: def weight_init(m):
    if any(isinstance(m, _m) for _m in [nn.Conv2d, nn.ConvTranspose2d]):
        nn.init.normal_(m.weight, mean=0.0, std=0.02)
    if m.bias is not None:
        nn.init.constant_(m.bias, 0)

```

1.6.4 Generator

It is the part of the layer that generates new images from the information.

```

[ ]: class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = self._init_encoder()
        self.decoder = self._init_decoder()

        self.out = nn.ConvTranspose2d(in_channels=128,
                                       out_channels=3,
                                       kernel_size=4,
                                       stride=2, padding=1)

        self.act = nn.Tanh()

        self.apply(weight_init)

    def forward(self, x):

        skips = []
        for layer in self.encoder:
            x = layer(x)
            skips.append(x)

        skips = reversed(skips[:-1])

        for layer, skip in zip(self.decoder, skips):
            x = layer(x)
            x = torch.cat((x, skip), dim=1)

        out = self.act(self.out(x))

        return out

```

```

def _init_encoder(self):
    kernel_size = 4
    in_channels = 64

    out_channels = [128, 256, 512, 512, 512, 512]

    encoder = [downsample(in_channels=3, out_channels=in_channels,
                           kernel_size=kernel_size, norm=False)]

    for out_ch in out_channels:
        encoder.append(downsample(in_channels=in_channels,
    ↪out_channels=out_ch,
                                   kernel_size=kernel_size))
        in_channels = out_ch

    return nn.Sequential(*encoder)

def _init_decoder(self):
    base_channels = 512
    kernel_size = 4
    in_channels = [512, 1024, 1024, 1024, 512, 256]

    out_channels = [512, 512, 512, 256, 128, 64]
    dropout = [True, True, True, False, False, False]

    decoder = [upsample(in_channels=in_ch, out_channels=out_ch,
                        kernel_size=kernel_size, dropout=drop)

    ↪for in_ch, out_ch, drop in zip(in_channels, out_channels,
    ↪dropout)]

    return nn.Sequential(*decoder)

```

```

[ ]: g = Generator()
x = torch.randn(1,3,256,256)

enc_out = g.encoder(x)
gen_out = g(x)

assert enc_out.shape[1] == 512
assert enc_out.shape[2] == x.shape[2] // 128
assert enc_out.shape[3] == x.shape[3] // 128

assert gen_out.shape[1] == x.shape[1]

```

```
assert gen_out.shape[2] == x.shape[2]
assert gen_out.shape[3] == x.shape[3]
```

```
[ ]: convolutions = []
for _, m in g.named_modules():
    if isinstance(m, nn.Conv2d):
        convolutions.append(m.weight.flatten().clone().detach().numpy())

convolutions = np.concatenate(convolutions)
convolutions.mean(), convolutions.std()
```

1.6.5 Discriminator

It is the part of the network which difference fake images from the real ones. In a cycle GAN it is necessary because we need to calculate the loss in some way and the discriminator will gives us this value from its evaluation.

```
[ ]: class Discriminator(nn.Module):
    def __init__(self):
        super().__init__()

        base_channels = 64
        kernel_size = 4
        self.discriminator = nn.Sequential(
            downsample(in_channels=3, out_channels=base_channels,
                        kernel_size=kernel_size, norm=False),
            downsample(in_channels=base_channels, out_channels=base_channels * 2,
                        kernel_size=kernel_size),
            downsample(in_channels=base_channels * 2, out_channels=base_channels * 4,
                        kernel_size=kernel_size),
            nn.ZeroPad2d(padding=(0, 2, 0, 2)),
            nn.Conv2d(in_channels=base_channels * 4, out_channels=base_channels * 8,
                        kernel_size=4, stride=1, bias=False),
            nn.GroupNorm(num_groups=base_channels * 8,
                        num_channels=base_channels * 8),
            nn.ZeroPad2d(padding=(0, 2, 0, 2)),
            nn.Conv2d(in_channels=base_channels * 8, out_channels=1,
                        kernel_size=4, stride=1),
        )

        self.apply(weight_init)

    def forward(self, x):
        return self.discriminator(x)
```

```
[ ]: d = Discriminator()
x = torch.randn(1,3,256,256)

assert d(x).shape[1] == 1

[ ]: torch_sample = torch.tensor(sample_photo, dtype=torch.float32)

torch_sample = (torch_sample / 255.0 - 0.5) / 0.5

torch_sample = torch_sample.permute(2,0,1)

torch_sample = torch_sample.unsqueeze(dim=0)
```

1.6.6 Testing the generator

From initial state, the generator will give some information about the Original Photo.

```
[ ]: monet_generator = Generator()

[ ]: with torch.inference_mode():
    monet_generator.eval()
    to_monet = monet_generator(torch_sample).detach()
    to_monet.shape

[ ]: plt.subplot(1,2,1)
plt.title("Original Photo")
plt.imshow(torch_sample[0].permute(1,2,0) * 0.5 + 0.5)

plt.subplot(1,2,2)
plt.title("Monet Photo")
plt.imshow(to_monet[0].permute(1,2,0) * 0.5 + 0.5)
```

1.6.7 Cycle GAN implementation

As I mention in the introduction of this section, a Cycle GAN distinguish the field of the desired target looping from one field to another. The basic logic of a Cycle GAN is two GANs competing between each other.

One trying to translate to field (Original -> Monet) and the other with (Monet -> Original).

In the training, both GANs generators are constantly trying to mislead their rivals.

Finally, it is needed to calculate the loss between each type of generated image: fake and original, from every field: normal and Monet. This will give the network a way to generate better images.

```
[ ]: class CycleGAN(nn.Module):
```

```

def __init__(self, lr=2e-4, lambda_cycle=10): # Not sure what that is yet
    super().__init__()

    self.gen_monet = Generator()
    self.optim_gen_monet = optim.Adam(self.gen_monet.parameters(),
                                      lr=lr,
                                      betas=(0.5, 0.999))

    self.gen_photo = Generator()
    self.optim_gen_photo = optim.Adam(self.gen_photo.parameters(),
                                      lr=lr,
                                      betas=(0.5, 0.999))

    self.disc_monet = Discriminator()
    self.optim_disc_monet = optim.Adam(self.disc_monet.parameters(),
                                       lr=lr,
                                       betas=(0.5, 0.999))

    self.disc_photo = Discriminator()
    self.optim_disc_photo = optim.Adam(self.disc_photo.parameters(),
                                       lr=lr,
                                       betas=(0.5, 0.999))

    self.loss_l2 = nn.MSELoss()
    self.loss_l1 = nn.L1Loss()

    self.lambda_cycle = lambda_cycle

def forward(self, x):
    return self.gen_monet(x)

def forward_step(self, x):
    x_monet, x_photo = x

    self.optim_zero_grad()

    x_fake_monet = self.gen_monet(x_photo)
    x_fake_photo = self.gen_photo(x_monet)

    d_real_monet = self.disc_monet(x_monet)
    d_fake_monet = self.disc_monet(x_fake_monet.detach())

    d_real_photo = self.disc_photo(x_photo)
    d_fake_photo = self.disc_photo(x_fake_photo.detach())

    d_loss_monet = self.disc_loss(x_real=d_real_monet,
                                x_fake=d_fake_monet)

    d_loss_photo = self.disc_loss(x_real=d_real_photo,

```



```

        x_fake=d_fake_photo)

    d_loss = d_loss_monet + d_loss_photo

    d_loss.backward()

    self.optim_disc_monet.step()
    self.optim_disc_photo.step()

    x_adv_monet = self.disc_monet(x_fake_monet)
    x_adv_photo = self.disc_photo(x_fake_photo)

    x_cycled_monet = self.gen_monet(x_fake_photo)
    x_cycled_photo = self.gen_photo(x_fake_monet)

    g_loss_adv_monet = self.loss_l2(torch.ones_like(x_adv_monet),
↪x_adv_monet)
    g_loss_adv_photo = self.loss_l2(torch.ones_like(x_adv_photo),
↪x_adv_photo)

    g_loss_cycle_monet = self.loss_l1(x_monet, x_cycled_monet) * self.
↪lambda_cycle
    g_loss_cycle_photo = self.loss_l1(x_photo, x_cycled_photo) * self.
↪lambda_cycle

    g_loss_monet = g_loss_adv_monet + g_loss_cycle_monet
    g_loss_photo = g_loss_adv_photo + g_loss_cycle_photo
    g_loss = g_loss_monet + g_loss_photo

    g_loss.backward()

    self.optim_gen_monet.step()
    self.optim_gen_photo.step()

    return {
        "g_monet" : g_loss_monet,
        "g_photo" : g_loss_photo,
        "d_monet" : d_loss_monet,
        "d_photo" : d_loss_photo
    }

def disc_loss(self, x_real, x_fake):
    loss_real = self.loss_l2(torch.ones_like(x_real), x_real)
    loss_fake = self.loss_l2(torch.zeros_like(x_fake), x_fake)

    return (loss_real + loss_fake) * 0.5

```

```
def optim_zero_grad(self):
    self.optim_gen_monet.zero_grad()
    self.optim_gen_photo.zero_grad()

    self.optim_disc_monet.zero_grad()
    self.optim_disc_photo.zero_grad()
```

```
[ ]: cycle_gan = CycleGAN()
```

```
[ ]: x_real = torch.ones(1, 3, 256, 256)
x_fake = torch.zeros_like(x_real)
cycle_gan.forward_step((x_real, x_fake))
```

```
[ ]: to_tensor = lambda x: torch.tensor(x, dtype=torch.float32).permute(2,0,1).
    ↪unsqueeze(0) / 127.5 - 1.0
from_tensor = lambda x: ((x.squeeze().permute(1,2,0).numpy() + 1.0) * 127.5).
    ↪astype(np.uint8)
```

```
[ ]: device = torch.device("cuda") if torch.cuda.is_available() else torch.
    ↪device("cpu")
```

```
[ ]: tensor_monet = to_tensor(sample_monet).to(device)
tensor_photo = to_tensor(sample_photo).to(device)

tensor_monet.min(), tensor_monet.max(), tensor_monet.shape, tensor_monet.device
```

1.7 Dataset Loading

Then I will load the full photos dataset.

```
[ ]: class MonetDataset(Dataset):
    def __init__(self, root_path):

        self.root_path = root_path
        self.photo_paths = os.listdir(f"{root_path}/photo_jpg")
        self.monet_paths = os.listdir(f"{root_path}/monet_jpg")

    def __len__(self):
        return len(self.photo_paths)

    def __getitem__(self, idx):
        path_photo = f"{root_path}/photo_jpg/{self.photo_paths[idx]}"

        path_monet = f"{root_path}/monet_jpg/{random.choices(self.monet_paths,
    ↪k=1)[0]}"

        x_photo = self.to_tensor(self.read_img(path_photo))
```

```

        x_monet = self.to_tensor(self.read_img(path_monet))

        return x_monet, x_photo

    def to_tensor(self, x):
        return torch.tensor(x, dtype=torch.float32).permute(2,0,1) / 127.5 - 1.0

    def read_img(self, path):
        return cv2.cvtColor(cv2.imread(path), cv2.COLOR_BGR2RGB)

```

1.8 Training

I will set the number of epochs to 50 and the batch size to 4. Originally the number of epochs was 30 and the batch 32. The main reason to this change is the closeness to desired performance and the 4GB max memory which, during training, it fills 2.4 GB of the total graphic memory.

```

[ ]: NUM_EPOCHS = 35
    BATCH_SIZE = 32

    device = torch.device("cuda") if torch.cuda.is_available() else torch.
        ↳device("cpu")

    loss_epoch = defaultdict(list)

[ ]: dataset = MonetDataset(root_path=root_path)
    data_loader = DataLoader(dataset=dataset, batch_size=BATCH_SIZE, shuffle=True,
        ↳pin_memory=True, num_workers=4)

    cycle_gan = CycleGAN(lr=1e-6).to(device)
    loss_epoch = defaultdict(list)

[ ]: VALIDATE EVERY = 0.1 * NUM_EPOCHS
    PHOTO_TO_MONET = []

    for i in range(1, NUM_EPOCHS + 1): # Starting at one for tqdm

        with tqdm(total=len(data_loader),
            desc=f"Epoch {i}/{NUM_EPOCHS}",
            ascii=True,
            colour="green") as pbar:

            loss_batch = defaultdict(list)
            for x in data_loader:
                x_monet = x[0].to(device)
                x_photo = x[1].to(device)

                loss = cycle_gan.forward_step((x_monet, x_photo))

```

```

        for k, v in loss.items():
            loss_batch[k].append(v.item())

        pbar.set_postfix({k:v.item() for k, v in loss.items()})
        pbar.update()

    for k, v in loss_batch.items():
        loss_epoch[k].append(np.mean(v))

    if i % VALIDATE_EVERY == 0:
        with torch.inference_mode(mode=True):
            cycle_gan.eval()
            monet_pred = cycle_gan(tensor_photo)

        PHOTO_TO_MONET.append(monet_pred)
        cycle_gan.train()

```

1.9 Loss Evaluation

In the following figure you will see the Generated images from both GANS and the real ones loss.

```

[ ]: plt.figure(figsize=(10,8))

plt.subplot(221)
plt.plot(range(len(loss_epoch["g_monet"])), loss_epoch["g_monet"])
plt.title("G Monet")

plt.subplot(222)
plt.plot(range(len(loss_epoch["g_photo"])), loss_epoch["g_photo"])
plt.title("G Photo")

plt.subplot(223)
plt.plot(range(len(loss_epoch["d_monet"])), loss_epoch["d_monet"])
plt.title("D Monet")

plt.subplot(224)
plt.plot(range(len(loss_epoch["d_photo"])), loss_epoch["d_photo"])
plt.title("D Photo")

plt.show()

```

[]:

1.10 Output evaluation

```
[ ]: plt.figure(figsize=(15,10))

plt.subplot(121)
plt.title("Photo")
plt.imshow(from_tensor(tensor_photo.cpu()))
plt.axis("off")

plt.subplot(122)
plt.title("Monet")
plt.imshow(sample_monet)
plt.axis("off")
```

```
[ ]: fig, axes = plt.subplots(2, 5, figsize=(20, 10))

axes = axes.flatten()

for i, img in enumerate(PHOTO_TO_MONET):
    axes[i].imshow(from_tensor(img.cpu()))
    axes[i].axis('off')
    axes[i].set_title(f"Photo to Monet at: {(i+1) * 10}%")

plt.tight_layout()
plt.show()
```

1.11 Submission saving

```
[ ]: import PIL
! mkdir ../images
```

```
[ ]: i = 1
data_loader = DataLoader(dataset=dataset, batch_size=1)
with torch.inference_mode(mode=True):
    cycle_gan.eval()

    for x in data_loader:
        photo, _ = x

        pred = cycle_gan(photo.to(device)).cpu().squeeze().numpy()
        pred = np.transpose(pred, (1,2,0))

        pred = (pred * 127.5 + 127.5).astype(np.uint8)

        pred_img = PIL.Image.fromarray(pred)
        pred_img.save("../images/" + str(i) + ".jpg")
```

```
i += 1
```

```
[ ]: import shutil  
shutil.make_archive("/kaggle/working/images", 'zip', "/kaggle/images")
```

1.12 Conclusion

In this work I trained a cycle GAN for Monet image generation, the overall project repository has two notebooks one with a local running and the other with an online Kaggle one. In future works it is possible to test other GPU architectures, batch sizes, and network architectures.